

INTRO to DATA SCIENCE

LECTURE 13: SUPPORT VECTOR MACHINES

I. DECISION TREES

II. FITTING DECISION TREES

III. OBJECTIVE FUNCTIONS

IV. REGULARIZATION

V. ENSEMBLE METHODS

BAGGING BOOSTING RANDOM FORESTS



Questions?

DATA EXPLORATION

SUPERVISED LEARNING: REGRESSION

SUPERVISED LEARNING: CLASSIFICATION

UNSUPERVISED LEARNING

VARIOUS TOPICS

LOGISTIC REGRESSION

NAIVE BAYES

RANDOM FORESTS

SUPPORT VECTOR MACHINES

COMPETITION

**Final outlines for your project
are due next lesson**

I. SUPPORT VECTOR MACHINES

II. REGULARIZATION

III. KERNELS

- **DESCRIBE WHAT THE SVM'S OBJECTIVE IS**
- **DESCRIBE THE EFFECT OF REGULARIZATION**
- **DESCRIBE WHAT KERNELS ARE**
- **APPLY SVMs IN SKLEARN**


- **DESCRIBE WHAT THE SVM'S OBJECTIVE IS**
- **DESCRIBE THE EFFECT OF REGULARIZATION**
- **DESCRIBE WHAT KERNELS ARE**
- **APPLY SVMS IN SKLEARN**

WE WON'T DIVE INTO THE MATHEMATICAL DETAILS TODAY
BUT THERE ARE LINKS IN THE REPO IF YOU'RE INTERESTED

I. SUPPORT VECTORS


*Suppose we have data with binary labels,
which we would like to classify*





*Suppose we have data with binary labels,
which we would like to classify*


*Recall that after fitting a classifier, we can draw
the **decision boundary** which separates the
two classes*



Suppose we have data with binary labels, which we would like to classify

*Recall that after fitting a classifier, we can draw the **decision boundary** which separates the two classes*


These boundaries depend on the specific optimization method. For example, logistic regression gives a different one than naive Bayes



Suppose we have data with binary labels, which we would like to classify

*Recall that after fitting a classifier, we can draw the **decision boundary** which separates the two classes*

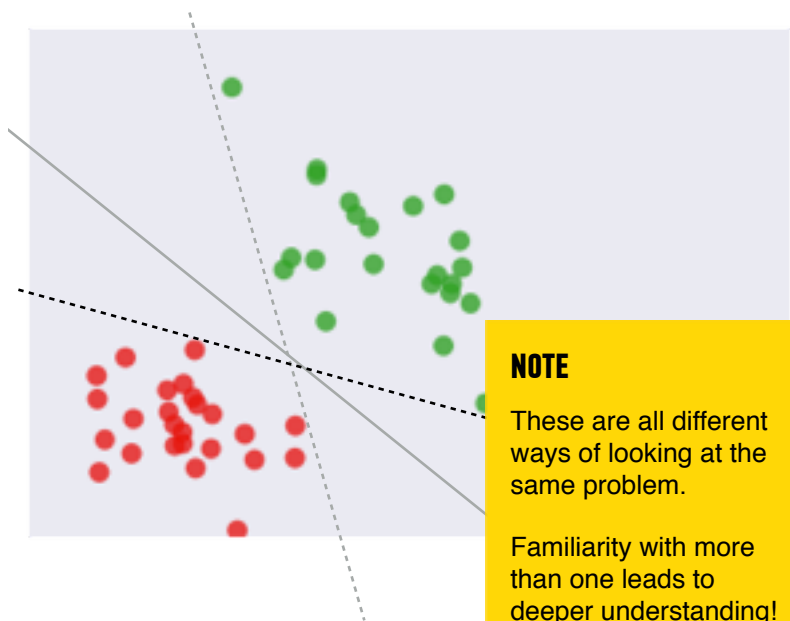
These boundaries depend on the specific optimization method. For example, logistic regression gives a different one than naive Bayes



Suppose we have data with binary labels, which we would like to classify

*Recall that after fitting a classifier, we can draw the **decision boundary** which separates the two classes*


These boundaries depend on the specific optimization method. For example, logistic regression gives a different one than naive Bayes




Suppose we have data with binary labels, which we would like to classify

*Recall that after fitting a classifier, we can draw the **decision boundary** which separates the two classes*

These boundaries depend on the specific optimization method. For example, logistic regression gives a different one than naive Bayes




Support Vector Machines is a classifier that explicitly constructs a decision boundary that geometrically “makes the most sense”



Support Vector Machines is a classifier that explicitly constructs a decision boundary that geometrically “makes the most sense”

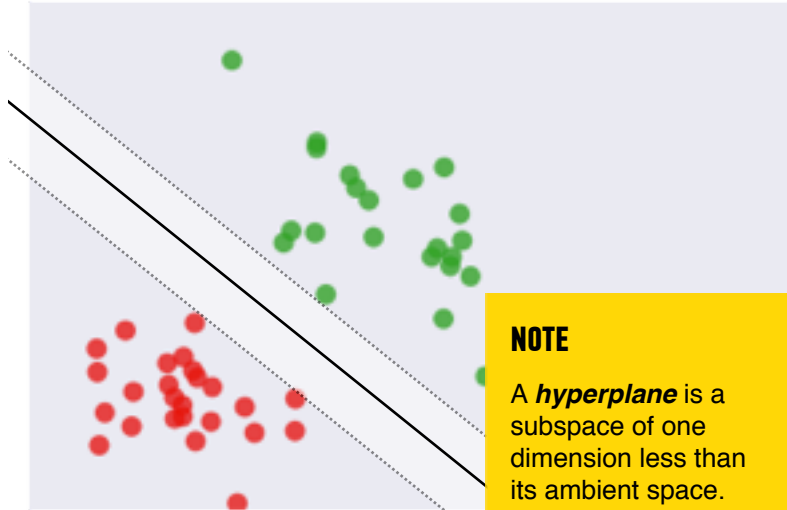
*The generalization error is equated with the geometric concept of **margin**, which is the region along the boundary that is free of points.*



Support Vector Machines is a classifier that explicitly constructs a decision boundary that geometrically “makes the most sense”

*The generalization error is equated with the geometric concept of **margin**, which is the region along the boundary that is free of points.*

*The goal of SVM is to create a linear decision boundary with the largest margin. This is called the **maximum margin hyperplane**.*

**NOTE**

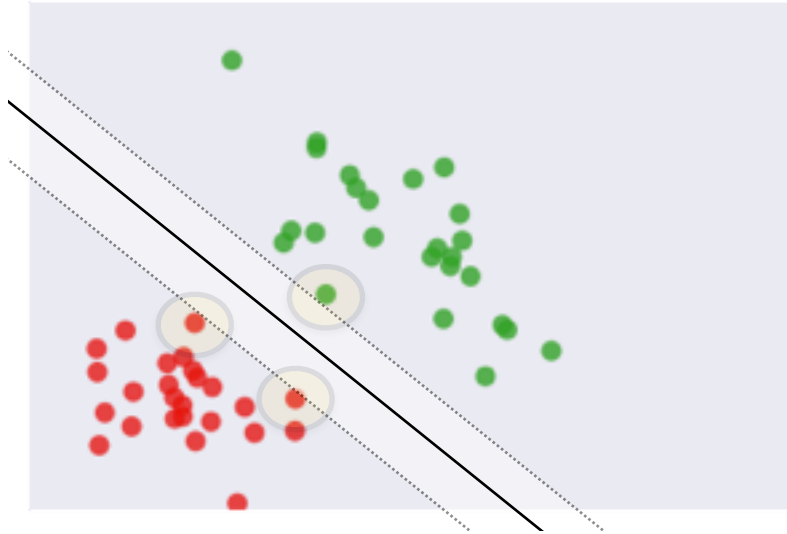
A **hyperplane** is a subspace of one dimension less than its ambient space.

In this 2D example, it is a line, and in a 3D space it is an ordinary plane.


Support Vector Machines is a classifier that explicitly constructs a decision boundary that geometrically “makes the most sense”

*The generalization error is equated with the geometric concept of **margin**, which is the region along the boundary that is free of points.*

*The goal of SVM is to create a linear decision boundary with the largest margin. This is called the **maximum margin hyperplane**.*




Notice that the margin depends only on a subset of the training data – the points nearest to the decision boundary.



Notice that the margin depends only on a subset of the training data – the points nearest to the decision boundary.


*These points are called the **support vectors**.*



Notice that the margin depends only on a subset of the training data – the points nearest to the decision boundary.

*These points are called the **support vectors**.*

The other points don't affect the construction of the hyperplane at all!



Notice that the margin depends only on a subset of the training data – the points nearest to the decision boundary.

*These points are called the **support vectors**.*

The other points don't affect the construction of the hyperplane at all!

*Finding the **maximum margin hyperplane** is a straightforward exercise in analytic geometry. (We won't go through the details here.)*

*Finding the **maximum margin hyperplane** is a straightforward exercise in analytic geometry. (We won't go through the details here.)*

*In short, this requires the optimization of a **convex** objective function.*



*Finding the **maximum margin hyperplane** is a straightforward exercise in analytic geometry. (We won't go through the details here.)*

*In short, this requires the optimization of a **convex** objective function.*

*Convex optimization are guaranteed to give **global optima**.*



So to summarize, what is a support vector machine?

So to summarize, what is a support vector machine?

An SVM is a binary linear classifier whose decision boundary is explicitly constructed to minimize generalization error.

recall:

binary classifier – *solves two-class problem*


linear classifier – *creates linear decision boundary*

II. REGULARIZATION


Let's apply a SVM to the dataset on the left



Let's apply a SVM to the dataset on the left.



*If the data are **linearly separable**, the training error is zero.*



Let's apply a SVM to the dataset on the left.

*If the data are **linearly separable**, the training error is zero.*

*The **margin** is nice and wide.*

Let's apply a SVM to the dataset on the left.

*If the data are **linearly separable**, the training error is zero.*

But what if our data has a single outlier?




Let's apply a SVM to the dataset on the left.

*If the data are **linearly separable**, the training error is zero.*

But what if our data has a single outlier?

*This will disproportionately impact the result, since the SVM tries to linearly separate **all** data.*




Let's apply a SVM to the dataset on the left.

*If the data are **linearly separable**, the training error is zero.*


But what if our data has a single outlier?

*This will disproportionately impact the result, since the SVM tries to linearly separate **all** data.*

*The **margin** is very small.*





Again, we'll need some sort of regularization.



Again, we'll need some sort of regularization.

*We want a **tradeoff** of the width of the margin and the number of misclassifications.*






Again, we'll need some sort of regularization.

*We want a **tradeoff** of the width of the margin and the number of misclassifications.*

We definitely need this when our data are not perfectly linearly separable.



Again, we'll need some sort of regularization.

*We want a **tradeoff** of the width of the margin and the number of misclassifications.*

We definitely need this when our data are not perfectly linearly separable.

*This is again a bias-variance tradeoff, which is done by **slack variables**.*

**NOTE**

We won't go into the mathematics this time, but in the course repo are several links to resources which will explain this if you're interested.

Again, we'll need some sort of regularization.

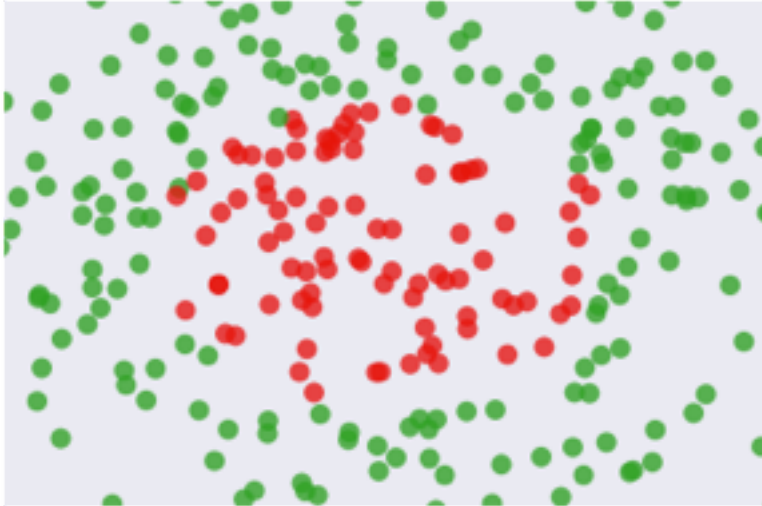
*We want a **tradeoff** of the width of the margin and the number of misclassifications.*

We definitely need this when our data are not perfectly linearly separable.


*This is again a bias-variance tradeoff, which is done by **slack variables**.*

III. KERNELS


What if our data is not linearly separable at all?



What if our data is not linearly separable at all?



Again, we could add polynomial features. (This might be computationally expensive.)




What if our data is not linearly separable at all?


Again, we could add polynomial features. (This might be computationally expensive.)

*We could also use **kernels**.*

*Add a **landmark** to the feature space.*



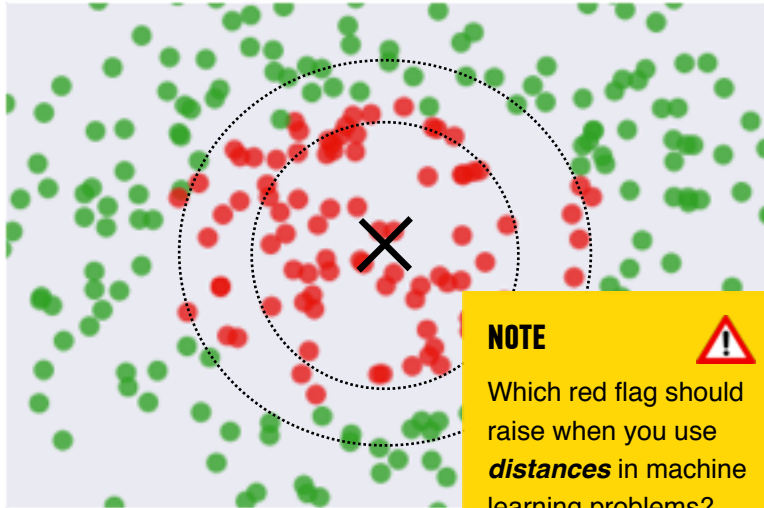
Add a **landmark** to the feature space.



For each point, compute the distance to this landmark: $\|x - l\|$

Add a **landmark** to the feature space.

For each point, compute the distance to this landmark: $\|x - l\|$



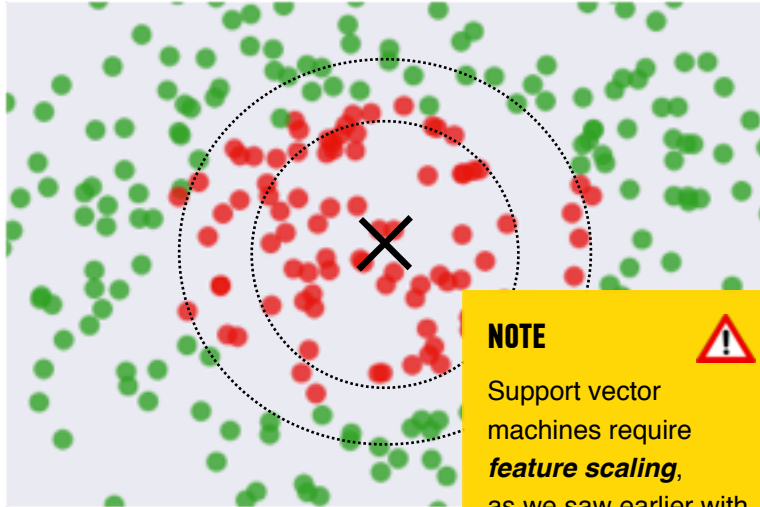
NOTE



Which red flag should raise when you use **distances** in machine learning problems?

Add a **landmark** to the feature space.

For each point, compute the distance to this landmark: $\|x - l\|$




NOTE



Support vector machines require **feature scaling**, as we saw earlier with k-nearest neighbors.


Add a **landmark** to the feature space.



For each point, compute the distance to this landmark: $\|x - l\|$

Then define the similarity as the **radius basis function (rbf)**

$$e^{-\frac{\|x-l\|^2}{2\sigma^2}}$$




Add a **landmark** to the feature space.

For each point, compute the distance to this landmark: $\|x - l\|$

Then define the similarity as the **radius basis function (rbf)**

$$e^{-\frac{\|x-l\|^2}{2\sigma^2}}$$

This is called a *Gaussian kernel*.




Add a **landmark** to the feature space.

For each point, compute the distance to this landmark: $\|x - l\|$

Then define the similarity as the **radius basis function (rbf)**

$$e^{-\frac{\|x-l\|^2}{2\sigma^2}}$$

This is called a *Gaussian kernel*.



*Add a **landmark** to the feature space.*


For each point, compute the distance to this landmark: $\|x - l\|$

*Then define the similarity as the **radius basis function (rbf)***


$$e^{-\frac{\|x-l\|^2}{2\sigma^2}}$$

This is called a Gaussian kernel.

How do we choose the right landmark?





How do we choose the right landmark?



How do we choose the right landmark?

Choose each sample as a landmark.






How do we choose the right landmark?

Choose each sample as a landmark.

Then replace all samples with the kernel values.




How do we choose the right landmark?

Choose each sample as a landmark.

Then replace all samples with the kernel values.

This will (non-linearly) transform your feature matrix \mathbf{X} from an $N \times n$ matrix to an $N \times N$ matrix.



How do we choose the right landmark?

Choose each sample as a landmark.

Then replace all samples with the kernel values.

This will (non-linearly) transform your feature matrix \mathbf{X} from an $N \times n$ matrix to an $N \times N$ matrix.

On the diagonal we have ones (each sample compared with itself).

*We can use a kernel to implicitly train our model in a high-dimensional feature space, **without** incurring additional computational complexity*

*We can use a kernel to implicitly train our model in a high-dimensional feature space, **without** incurring additional computational complexity*

As long as the kernel function satisfies certain conditions, we can perform the same methods for the maximum margin hyperplane

*We can use a kernel to implicitly train our model in a high-dimensional feature space, **without** incurring additional computational complexity*

As long as the kernel function satisfies certain conditions, we can perform the same methods for the maximum margin hyperplane

Nonlinear classification is then obtained by creating a linear decision boundary in the higher-dimensional space

*We can use a kernel to implicitly train our model in a high-dimensional feature space, **without** incurring additional computational complexity*

As long as the kernel function satisfies certain conditions, we can perform the same methods for the maximum margin hyperplane

Nonlinear classification is then obtained by creating a linear decision boundary in the higher-dimensional space

NOTE

These conditions are contained in a result called *Mercer's theorem*.

Kernels can be used as preprocessing step at other models as well.

Kernels can be used as preprocessing step at other models as well.

However, they tend to be very computationally expensive.

Kernels can be used as preprocessing step at other models as well.

However, they tend to be very computationally expensive.

The SVM is far more efficient, so using kernels is more practical.


Some popular kernels


linear kernel $k(\mathbf{x}, \mathbf{x}') = \langle \mathbf{x}, \mathbf{x}' \rangle$





polynomial kernel $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + 1)^d$

Gaussian kernel $k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$

The hyperparameters d and γ affect the flexibility of the dec. boundary







*SVMs (and **kernel methods** in general) are versatile, powerful, and popular techniques that can produce accurate results for a wide array of classification problems.*

*SVMs (and **kernel methods** in general) are versatile, powerful, and popular techniques that can produce accurate results for a wide array of classification problems.*

The main disadvantage of SVMs is the lack of intuition they produce.

These models are truly black boxes!

INTRO TO DATA SCIENCE

DISCUSSION